

dm510 project 3

Danny Jensen
danje14

Group:
Lea Fog-Fredsgaard
Michelle Dung Hoang

April 2018

1 Introduction

In this project a small driver is to be implemented, using the *scull pipe.c* as a reference. since much of the code is basically copy paste from *scull pipe.c* this rapport will focus on the different places.

2 Design

As *pipe.c* already has a working device it is used as a base and then simply split the struct up into a device struct and a buffer struct. For example the pointers to the places that can be read and wrote from was first designed to be in the devices, but then a problem would occur because device 0 would have a write pointer to buffer 1 buffer and a read pointer to buffer 0 buffer. But then a problem would occur because it was not possible to compare these two pointers because device 0 and device 1 could not talk together. So the pointers was then to be placed in the buffers instead, which made it all easier.

Now that the structs are designed, the init has no real design choices everything just needs to be initialized. The only design choices is what needs to be returned if *kmalloc* fails, and that is that there is not enough memory. And if *cdev_add* fails it just says try again, there is a very small chance it fails so if it does the process should just try again.

In the open function there needs to be a check that says if there are to many processes trying to open, then only the number of processes that was defined can have access and the rest gets an error saying tat they need to try again. the most important part here is that there is at least one designated spot for a reader, so that there wont be a deadlock where all the processes are readers that just waits on input but there are no space for a writer.

In the read function there needs to be a way that moves the read pointer in the buffer and the way it needs to move it is by simply moving it the number of bytes that was read.

The write function needs to do the same as the read, just with the write pointer in the buffer.

3 Implementation

First off the buffer struct and the device struct is made splitting most of the *pipe.c* device struct up. The buffer gets the wait ques because it makes sense that it is the buffer that wakes up the next process that can either write or read to it. The buffer also needs some pointer for the start and the end of the buffer. The buffer also keeps track of how many processes are interacting with it, the numbers are stored in *nreaders* and *nwriters*. The last thing that the buffer have are read and write pointers, to point where the write function can start writing and the read function can start reading.

```

1 struct dm510_buffer {
2     wait_queue_head_t inq, outq; /* read and write queues */
3     char *buffer, *end;          /* begin of buf, end of buf */
4     int nreaders, nwriters;      /* number of openings for r/w */
5     char *buffer_rp, *buffer_wp; /* where to read, where to write */
6 };

```

Now the device struct needs pointers to the buffers it writes from and reads from. It also holds the *cdev* struct. And of course a mutex lock which locks the device so it can't be interrupted.

```

1 struct dm510_dev {
2     struct dm510_buffer *read_buffer; /* pointer to buffer for
3     reading */
4     struct dm510_buffer *write_buffer; /* pointer to buffer for
5     writing */
6     struct cdev cdev;                  /* Char device structure */
7     struct mutex mutex;                /* mutual exclusion semaphore
8     */
9 };

```

Now in the init function *kmalloc* is used to allocate space in the kernel for both devices and the buffers and the buffer space. Of course *kmalloc* can fail and then an error is returned, and if anything had successfully gotten allocated some space it needs to be freed so as not to get dead parts of the memory.

```

1 dm510_dev_1 = kmalloc(sizeof(struct dm510_dev), GFP_KERNEL);
2 if (!dm510_dev_1) {
3     kfree(dm510_dev_0);
4     return -ENOMEM;
5 }

```

Then all the integers are set to the right starting value and pointers are set to the right places. (This was actually something we forgot to do with the read and write pointers and it took us almost an entire day to figure that mistake out. Lesson learned :))

```

1 dm510_buffer_0->nwriters = 0;
2 dm510_buffer_0->buffer_rp = dm510_buffer_0->buffer;

```

The wait queues, the mutex locks and the *cdev* needs to be initialized for later use to.

```

1 init_waitqueue_head(&dm510_buffer_0->inq);
2 mutex_init(&dm510_dev_0->mutex);
3
4 cdev_init(&(dm510_dev_0->cdev), &dm510_fops);

```

Of course the owner in *cdev* also needs to be set.

```

1 dm510_dev_0->cdev.owner = THIS_MODULE;

```

The release function must undo everything the init function has created, and it does so in reverse order, starting with the last thing init did. It is done in reverse so as to not delete any pointers that should be used to find

```

1 if (!dev_holder){
2     return;
3 }
4 /* cleans up in reverse order */
5 cdev_del(&dm510_dev_1->cdev); /* removes devices from the dev_t
   holder */
6 cdev_del(&dm510_dev_0->cdev);
7 kfree(dm510_buffer_1->buffer); /*free content of buffer*/
8 kfree(dm510_buffer_0->buffer);
9 kfree(dm510_buffer_1); /*free buffer*/
10 kfree(dm510_buffer_0);
11 kfree(dm510_dev_1); /*free device*/
12 kfree(dm510_dev_0);
13
14 /* removes the dev_t place holder */
15 unregister_chrdev_region(dev_holder, DEVICECOUNT);

```

Now in the open function there is a check build in that ensures that the number of processes dont get larger that specified and also that a deadlock doesn't happen where there is max number of processes running and they all are readers. This would mean that all the processes sleep an one process waits for something to be written but there is no space for a writer. This check is a if statement where if the max number of processes is running the new process gets an error stating to try again, and hopefully by the time the processes reties some other process got done and a space opened up.

```

1 if (dev->read_buffer->nreaders >= number_proc-1 && (filp->f_mode &
   FMODE_READ)) {
2     mutex_unlock(&dev->mutex);
3     return -EAGAIN;
4 } else if (dev->read_buffer->nreaders + dev->write_buffer->nwriters
   >= number_proc){ /* if there is a writer and max # of processes
   is running */
5     mutex_unlock(&dev->mutex);
6     return -EAGAIN;

```

now if theres already a writing processes running the new writing processes is put to sleep and the old writing process will wake a new writing process when done.

```

1 dev->write_buffer->nwriters++;
2 /* if there are one writer allready the rest is put to sleep */
3 if(wait_event_interruptible(dev->write_buffer->inq, (dev->
   write_buffer->nwriters >= 1))){
4     mutex_unlock(&dev->mutex);
5     return -ERESTARTSYS;
6 }

```

In read an *access_ok* is done to ensure the user space pointer is valid, if not an error is returned and the processes stopped.

```

1 /* if access is not ok return error */
2 if (!access_ok(VERIFY_WRITE, buf, count)){
3     mutex_unlock (&dev->mutex);
4     return -EACCES;
5 }

```

Now if the user space pointer is ok, the *copy_to_user* gets the requested number of bytes and puts them into user space. Then the read pointer is moved to the right place and if the right place was at the end of the buffer the read pointer is moved to the start of the buffer again.

```

1  /* copies to user and the remaining # of bytes not copied is stored
   in remaining */
2  remaining = copy_to_user(buf, dev->read_buffer->buffer_rp, count);
3
4  /* moving the read pointer to the next non read data */
5  dev->read_buffer->buffer_rp += count;
6
7  /* move read pointer from end to start */
8  if (dev->read_buffer->buffer_rp >= dev->read_buffer->end){
9      dev->read_buffer->buffer_rp = dev->read_buffer->buffer; /*
   wrapped */
10 }

```

The write function does the exact same thing just instead of *copy_to_user* it uses *copy_from_user* and the int variable *remaining* is used to subtract from the count to ensure that a read process does not read garbage.

```

1  dev->write_buffer->buffer_wp += count - remaining;

```

Also the *remaining* variable is also used to subtract from the count in the return value so as to know how much actually was written to the buffer.

```

1  return count - remaining; //return number of bytes written

```

Now for the IO control, the two cases are used to either change the buffer size or to change the max number of processes.

The first case *IOC_RESETBUFFER* first ensures that the argument handed to it is not 0 or negative, because it would be impossible to allocate negative space and it makes no sense trying to change the buffer size to 0.

If the argument is larger than 0, the two buffers are freed and two new spaces are allocated for the buffers, then the different pointer like end and the read and write pointers are reset, if the read and write pointers are not reset they would point to the old places in memory and that would be a problem, it would mean that a process would write into another process' memory and maybe destroy that process. And if a process read from the old pointer it would only get garbage.

Now that the buffer sizes has been changed and the pointer set at the right places the new buffer size is returned.

```

1  case IOC_RESETBUFFER:
2      buffer_size = arg;
3      if (buffer_size <= 0) {
4          return -EINVAL;
5      }
6      kfree(dm510_buffer_1->buffer);
7      kfree(dm510_buffer_0->buffer);
8      dm510_buffer_0->buffer = kmalloc(sizeof(char*) * buffer_size,
   GFP_KERNEL);
9      if (!dm510_buffer_0->buffer){
10         return -ENOMEM;

```

```

11 }
12 dm510_buffer_1->buffer = kmalloc(sizeof(char*) *buffer_size ,
    GFP_KERNEL);
13 if (!dm510_buffer_1->buffer){
14     return -ENOMEM;
15 }
16
17 dm510_buffer_0->end = dm510_buffer_0->buffer + buffer_size;
18 dm510_buffer_1->end = dm510_buffer_1->buffer + buffer_size;
19
20 dm510_buffer_0->buffer_rp = dm510_buffer_0->buffer;
21 dm510_buffer_0->buffer_wp = dm510_buffer_0->buffer;
22
23 dm510_buffer_1->buffer_rp = dm510_buffer_1->buffer;
24 dm510_buffer_1->buffer_wp = dm510_buffer_1->buffer;
25
26 return buffer_size;

```

IOC_RESETPROC resets the max number of processes and this one is quite easy, it simply assures that there can be at least 2 processes so that it is assured that a writer and a reader can be at the same time. If this check was set to at least 1 then if the max processes was set to 1 there would be no space for a read process because of the check in the open function.

Now if the argument passed onto *IOC_RESETPROC* is valid, *number_proc* is simply changed to the new value

```

1 case IOC_RESETPROC:
2     if (arg <=1) {
3         return -EINVAL;
4     }
5     number_proc = arg;
6     return number_proc;

```

4 Test

The test was video captured so it will be refereed to with time stamps here instead of showing pictures.

For the first test is to show that *dm510_load* works, this is shown at 0.36 in the video, and it loads just fine.

The next test is then *moduletest.c*, as shown at 0.41, this test prints out 2 tests and what the expected result should be in these 2 tests. And as shown the first test returns -930547960 as it should. The second returned -830547960, this is also equal to the expected value so from this it is concluded that *moduletest* was passed successfully.

Then just to be sure the devices is unloaded and loaded back again, this is done so the next test is started with everything at the starting points.

Now to test the *ioctl*, the first test is to change te buffer size, this happens at 0.55 in video. The new buffer size is set to 500 and the new buffer size is printed out just to be sure that it was changed.

The last *ioctl* test is to test if max number of processes can be changed, this

happens at 1.00. And just like the buffer size change the max number of processes value is also printed out so that it is clear that it was changed.

Now since the buffer size and max number of processes got tampered with the devices are unloaded and loaded back again for the last test.

The last test is started at 1.12 and the first test is to write to device 0 and read the just inserted text from device 1, the result is written out at 1.14. A simple "hello" string was written and read back to the user which is a success. The buffer size was changed to 15 for the last test.

Now device 1 is written to and device 0 is read from, this happen at 1.15. This time the string was "world" just to show that it is a new string and not the old that was returned again.

So now both devices was checked with both write and read functions and passed this test to.

The last test is to show that the write function wont write over data that has not yet been read. This test is at 1.16 in the video. As shown the writer tries to insert 15 bytes into the buffer, but since the previous test inserted 5 already (which has been read to), there is only space for 10 bytes until the write pointer hits the end of the buffer. The next writer is also trying to insert 15 but the read pointer is standing at 5 so it can only insert 4 so as not to get the write and read pointer standing in the same place. If that would happen it would look like there is nothing to read in the buffer.

5 Discussion

If multiple processes tried to use the devices at the same time they would be put to sleep using *wait_event_interruptible* that one process can do what it needs and when this process is done it should wake another process up from its sleep using *wake_up_interruptible*.

Now if to many processes tries use the device and the max number of processes is reached every other process get an error code returned stating it should try again.

6 Conclusion

All in all this project was successful and fun to work with, even though there where many times where it felt like it was impossible to finish because of some little bug which was impossible to find.

Thanks to this project a new found respect was found for the ones who work with drivers and suc every day.

7 Appendix

7.1 dm510_dev.c

```
1
2 /* Prototype module for second mandatory DM510 assignment */
3 #ifndef __KERNEL__
4 # define __KERNEL__
5 #endif
6 #ifndef MODULE
7 # define MODULE
8 #endif
9
10 #include "ioctl.h"
11 #include <linux/sched.h>
12 #include <linux/sched/signal.h>
13
14 #include <linux/module.h>
15 #include <linux/init.h>
16 #include <linux/slab.h>
17 #include <linux/kernel.h>
18 #include <linux/fs.h>
19 #include <linux/errno.h>
20 #include <linux/types.h>
21 #include <linux/wait.h>
22 // #include <asm/uaccess.h>
23 #include <linux/uaccess.h>
24 #include <linux/semaphore.h>
25 // #include <asm/system.h>
26 #include <asm/switch_to.h>
27 #include <linux/cdev.h>
28
29
30 /* Prototypes — this would normally go in a .h file */
31 static int dm510_open( struct inode*, struct file* );
32 static int dm510_release( struct inode*, struct file* );
33 static ssize_t dm510_read( struct file*, char*, size_t, loff_t* );
34 static ssize_t dm510_write( struct file*, const char*, size_t,
35                             loff_t* );
36 long dm510_ioctl( struct file *filp, unsigned int cmd, unsigned long
37                  arg );
38
39 #define DEVICE_NAME "dm510_dev" /* Dev name as it appears in /proc/
40                                devices */
41 #define MAJOR_NUMBER 254
42 #define MIN_MINOR_NUMBER 0
43 #define MAX_MINOR_NUMBER 1
44
45 #define DEVICE_COUNT 2
46 /*
47 #define DM510_IOC_MAGIC 9
48
49 #define IOC_RESETBUFFER _IO(DM510_IOC_MAGIC, 0)
50 #define IOC_RESETPROC _IO(DM510_IOC_MAGIC, 1)
51
52 #define IOC_NUMCASES 2*/
53 /* end of what really should have been in a .h file */
```



```

51
52 /* file operations struct */
53 static struct file_operations dm510_fops = {
54     .owner    = THIS_MODULE,
55     .read     = dm510_read,
56     .write    = dm510_write,
57     .open     = dm510_open,
58     .release  = dm510_release,
59     .unlocked_ioctl = dm510_ioctl
60 };
61
62 struct dm510_buffer {
63     wait_queue_head_t inq, outq;      /* read and write queues */
64     char *buffer, *end;               /* begin of buf, end of buf */
65     int nreaders, nwriters;           /* number of openings for r/w */
66     /*
67     char *buffer_rp, *buffer_wp;      /* where to read, where to
68                                         write */
69 };
70
71 struct dm510_dev {
72     struct dm510_buffer *read_buffer; /*pointer to buffer for
73                                         reading*/
74     struct dm510_buffer *write_buffer; /*pointer to buffer for
75                                         writing*/
76     struct cdev cdev;                 /* Char device structure */
77     struct mutex mutex;               /* mutual exclusion semaphore */
78     /*
79 };
80
81 /*initialising*/
82 struct dm510_buffer *dm510_buffer_0;
83 struct dm510_buffer *dm510_buffer_1;
84
85 struct dm510_dev *dm510_dev_0;
86 struct dm510_dev *dm510_dev_1;
87
88 int buffer_size = 3000;
89 int number_proc=10;
90
91 /*type that is defined and is used to hold device numbers (major,
92     minor)*/
93 dev_t dev_holder;
94 /* called when module is loaded */
95 int dm510_init_module( void ) {
96
97     /* initialization code belongs here */
98     int err;
99     /*creates the first device*/
100    dev_holder = MKDEV(MAJOR_NUMBER, MIN_MINOR_NUMBER);
101    /*registration takes a pointer and a name*/
102    err = register_chrdev_region(dev_holder, DEVICE_COUNT, DEVICE_NAME);
103    ;
104    if(err != 0){
105        printk(KERN_NOTICE "Unable to get region, error %d\n", err);
106        return -ENODEV;
107    }
108 }

```

```

101
102 /* allocating space for devices and buffers
103  * and freeing already allocated space if error occurs */
104 dm510_dev_0 = kmalloc(sizeof(struct dm510_dev), GFP_KERNEL);
105 if (!dm510_dev_0){
106     return -ENOMEM;
107 }
108 dm510_dev_1 = kmalloc(sizeof(struct dm510_dev), GFP_KERNEL);
109 if (!dm510_dev_1){
110     kfree(dm510_dev_0);
111     return -ENOMEM;
112 }
113 dm510_buffer_0 = kmalloc(sizeof(struct dm510_buffer), GFP_KERNEL)
114 ;
115 if (!dm510_buffer_0){
116     kfree(dm510_dev_0);
117     kfree(dm510_dev_1);
118     return -ENOMEM;
119 }
120 dm510_buffer_1 = kmalloc(sizeof(struct dm510_buffer), GFP_KERNEL)
121 ;
122 if (!dm510_buffer_1){
123     kfree(dm510_dev_0);
124     kfree(dm510_dev_1);
125     kfree(dm510_buffer_0);
126     return -ENOMEM;
127 }
128 /* allocating space for text in buffer */
129 dm510_buffer_0->buffer = kmalloc(sizeof(char*) *buffer_size ,
130 GFP_KERNEL);
131 if (!dm510_buffer_0->buffer){
132     kfree(dm510_dev_0);
133     kfree(dm510_dev_1);
134     kfree(dm510_buffer_0);
135     kfree(dm510_buffer_1);
136     return -ENOMEM;
137 }
138 dm510_buffer_1->buffer = kmalloc(sizeof(char*) *buffer_size ,
139 GFP_KERNEL);
140 if (!dm510_buffer_1->buffer){
141     kfree(dm510_dev_0);
142     kfree(dm510_dev_1);
143     kfree(dm510_buffer_0->buffer);
144     kfree(dm510_buffer_0);
145     kfree(dm510_buffer_1);
146     return -ENOMEM;
147 }
148 /* initialising all int's in the buffer to the right size */
149 dm510_buffer_0->nreaders = 0;
150 dm510_buffer_0->nwriters = 0;
151 dm510_buffer_0->buffer_rp = dm510_buffer_0->buffer;
152 dm510_buffer_0->buffer_wp = dm510_buffer_0->buffer;
153
154 dm510_buffer_1->nreaders = 0;
155 dm510_buffer_1->nwriters = 0;
156 dm510_buffer_1->buffer_rp = dm510_buffer_1->buffer;
157 dm510_buffer_1->buffer_wp = dm510_buffer_1->buffer;

```

```

154
155     dm510_buffer_0->end = dm510_buffer_0->buffer + buffer_size;
156     dm510_buffer_1->end = dm510_buffer_1->buffer + buffer_size;
157
158     /* initialize read and write queues */
159     init_waitqueue_head(&dm510_buffer_0->inq);
160     init_waitqueue_head(&dm510_buffer_0->outq);
161     init_waitqueue_head(&dm510_buffer_1->inq);
162     init_waitqueue_head(&dm510_buffer_1->outq);
163
164     /* initialise the mutex locks */
165     mutex_init(&dm510_dev_0->mutex);
166     mutex_init(&dm510_dev_1->mutex);
167
168     /* initialize a cdev structure */
169     cdev_init(&(dm510_dev_0->cdev), &dm510_fops);
170     cdev_init(&(dm510_dev_1->cdev), &dm510_fops);
171
172     /* setting the owner */
173     dm510_dev_0->cdev.owner = THIS_MODULE;
174     dm510_dev_1->cdev.owner = THIS_MODULE;
175
176     /* setting the right read and write buffers in the devices */
177     dm510_dev_0->read_buffer = dm510_buffer_0;
178     dm510_dev_0->write_buffer = dm510_buffer_1;
179     dm510_dev_1->read_buffer = dm510_buffer_1;
180     dm510_dev_1->write_buffer = dm510_buffer_0;
181
182     /* add the devices to the dev_t place holder */
183     err = cdev_add(&dm510_dev_0->cdev, dev_holder, 1);
184     if (err != 0) {
185         printk(KERN_NOTICE "Error %d adding cdev", err);
186         return -EAGAIN;
187     }
188
189     err = cdev_add(&dm510_dev_1->cdev, dev_holder+1, 1);
190     if (err != 0) {
191         printk(KERN_NOTICE "Error %d adding cdev", err);
192         return -EAGAIN;
193     }
194
195     printk(KERN_INFO "DM510: Hello from your device!\n");
196     return 0;
197 }
198
199 /* Called when module is unloaded */
200 void dm510_cleanup_module( void ) {
201     /* clean up code belongs here */
202
203     if (!dev_holder){
204         return;
205     }
206     /* cleans up in reverse order */
207     cdev_del(&dm510_dev_1->cdev); /* removes devices from the dev_t
        holder */
208     cdev_del(&dm510_dev_0->cdev);
209     kfree(dm510_buffer_1->buffer); /*free content of buffer*/

```

```

210 kfree(dm510_buffer_0->buffer);
211 kfree(dm510_buffer_1);           /*free buffer*/
212 kfree(dm510_buffer_0);
213 kfree(dm510_dev_1);             /*free device*/
214 kfree(dm510_dev_0);
215
216 /* removes the dev_t place holder */
217 unregister_chrdev_region(dev_holder, DEVICE_COUNT);
218
219 printk(KERN_INFO "DM510: Module unloaded.\n");
220 }
221
222 /* Called when a process tries to open the device file */
223 static int dm510_open( struct inode *inode, struct file *filp ) {
224     /* device claiming code belongs here */
225     struct dm510_dev *dev;
226
227     /* puts the devices into filp */
228     dev = container_of(inode->i_cdev, struct dm510_dev, cdev);
229     filp->private_data = dev;
230
231     /* locks the process */
232     if (mutex_lock_interruptible(&dev->mutex)){
233         return -ERESTARTSYS;
234     }
235
236     /* make checks to ensure that number of processes is kept */
237     if (dev->read_buffer->nreaders >= number_proc-1 && ( filp->f_mode
238         & FMODE_READ)) {
239         mutex_unlock(&dev->mutex);
240         return -EAGAIN;
241     } else if (dev->read_buffer->nreaders + dev->write_buffer->
242         nwriters >= number_proc){ /* if there is a writer and max # of
243         processes is running */
244         mutex_unlock(&dev->mutex);
245         return -EAGAIN;
246     } else { /* if its either a write, or there are fewer than max
247         processes */
248         if (filp->f_mode & FMODE_READ){
249             dev->read_buffer->nreaders++;
250         }
251         if (filp->f_mode & FMODE_WRITE){
252             if (filp->f_flags & O_NONBLOCK) {
253                 mutex_unlock(&dev->mutex);
254                 return -EAGAIN;
255             }
256             dev->write_buffer->nwriters++;
257             /* if there are one writer already the rest is put to sleep
258             */
259             if (wait_event_interruptible(dev->write_buffer->inq, (dev->
260                 write_buffer->nwriters >= 1))){
261                 mutex_unlock(&dev->mutex);
262                 return -ERESTARTSYS;
263             }
264         }
265     }
266     mutex_unlock(&dev->mutex);

```

```

261     return 0;
262 }
263
264 /* Called when a process closes the device file. */
265 static int dm510_release( struct inode *inode, struct file *filp )
266 {
267     /* device release code belongs here */
268
269     struct dm510_dev *dev = filp->private_data;
270
271     mutex_lock(&dev->mutex);
272     if ( filp->f_mode & FMODE_READ){
273         dev->read_buffer->nreaders--;
274     }
275     if ( filp->f_mode & FMODE_WRITE){
276         dev->write_buffer->nwriters--;
277     }
278     mutex_unlock(&dev->mutex);
279
280     return 0;
281 }
282
283 /* Called when a process, which already opened the dev file,
284    attempts to read from it. */
285 static ssize_t dm510_read( struct file *filp,
286     char *buf, /* The buffer to fill with data */
287     size_t count, /* The max number of bytes to read */
288     loff_t *f_pos ) /* The offset in the file */
289 {
290     /* read code belongs here */
291
292     struct dm510_dev *dev = filp->private_data;
293
294     int remaining;
295
296     if ( mutex_lock_interruptible(&dev->mutex)){
297         return -ERESTARTSYS;
298     }
299     while (dev->read_buffer->buffer_rp == dev->read_buffer->buffer_wp
300 ) { /* nothing to read */
301         mutex_unlock(&dev->mutex); /* release the lock */
302         if ( filp->f_flags & O_NONBLOCK){
303             return -EAGAIN;
304         }
305         if ( wait_event_interruptible(dev->read_buffer->inq, (dev->
306 read_buffer->buffer_rp != dev->read_buffer->buffer_wp))){
307             return -ERESTARTSYS; /* signal: tell the fs layer to handle
308 it */
309         }
310         /* otherwise loop, but first reacquire the lock */
311         if ( mutex_lock_interruptible(&dev->mutex)){
312             return -ERESTARTSYS;
313         }
314     }
315
316     /* ok, data is there, return something */

```

```

313 if (dev->read_buffer->buffer_wp > dev->read_buffer->buffer_rp){
314     count = min(count, (size_t)(dev->read_buffer->buffer_wp - dev->
        read_buffer->buffer_rp));
315 } else { /* the write pointer has wrapped, return data up to dev->
        end */
316     count = min(count, (size_t)(dev->read_buffer->end - dev->
        read_buffer->buffer_rp));
317 }
318
319 /* if access is not ok return error */
320 if (!access_ok(VERIFY_WRITE, buf, count)){
321     mutex_unlock (&dev->mutex);
322     return -EACCES;
323 }
324
325 /* copies to user and the remaining # of bytes not copied is
        stored in remaining */
326 remaining = copy_to_user(buf, dev->read_buffer->buffer_rp, count)
        ;
327
328 /* moving the read pointer to the next non read data */
329 dev->read_buffer->buffer_rp += count;
330
331 /* move read pointer from end to start */
332 if (dev->read_buffer->buffer_rp >= dev->read_buffer->end){
333     dev->read_buffer->buffer_rp = dev->read_buffer->buffer; /*
        wrapped */
334 }
335 mutex_unlock (&dev->mutex);
336
337 /* finally , awake any writers and return */
338 wake_up_interruptible(&dev->read_buffer->outq);
339
340 /* return number of bytes read */
341 return count;
342 }
343
344 /* Called when a process writes to dev file */
345 static ssize_t dm510_write( struct file *filp ,
346     const char *buf, /* The buffer to get data from */
347     size_t count, /* The max number of bytes to write */
348     loff_t *f_pos ) /* The offset in the file */
349 {
350     /* write code belongs here */
351     struct dm510_dev *dev = filp->private_data;
352     int remaining;
353
354     /* if trying to read bytes under 1 */
355     if (count < 1)
356     {
357         return -EINVAL;
358     }
359
360     if (mutex_lock_interruptible(&dev->mutex)){
361         return -ERESTARTSYS;
362     }
363

```

```

364     if (dev->write_buffer->buffer_wp >= dev->write_buffer->buffer_rp)
365     {
366         count = min(count, (size_t)(dev->write_buffer->end - dev->
367         write_buffer->buffer_wp)); /* to end-of-buf */
368     } else { /* the write pointer has wrapped, fill up to rp-1 */
369         count = min(count, (size_t)(dev->write_buffer->buffer_rp - dev
370         ->write_buffer->buffer_wp - 1));
371     }
372
373     if (!access_ok(VERIFY_WRITE, buf, count)){
374         mutex_unlock (&dev->mutex);
375         return -EACCES;
376     }
377
378     remaining = copy_from_user((dev->write_buffer->buffer_wp), buf,
379     count);
380
381     dev->write_buffer->buffer_wp += count - remaining;
382
383     if (dev->write_buffer->buffer_wp >= dev->write_buffer->end){
384         dev->write_buffer->buffer_wp = dev->write_buffer->buffer; /*
385         wrapped */
386     }
387     mutex_unlock(&dev->mutex);
388     //dev->write_buffer->buffer_wp = dev->write_p;
389
390     wake_up_interruptible(&dev->write_buffer->inq); /* blocked in
391     read() and select() */
392
393     return count - remaining; //return number of bytes written
394 }
395
396 /* called by system call icotl */
397 long dm510_ioctl(
398     struct file *filp,
399     unsigned int cmd, /* command passed from the user */
400     unsigned long arg ) /* argument of the command */
401 {
402     /* ioctl code belongs here */
403     printk(KERN_INFO "DM510: ioctl called.\n");
404
405     if (_IOC_TYPE(cmd) != DM510_IOC_MAGIC){
406         return -ENOTTY;
407     }
408     if (_IOC_NR(cmd) > IOC_NUMCASES){
409         return -ENOTTY;
410     }
411
412     switch(cmd){
413     case IOC_RESETBUFFER:
414         buffer_size = arg;
415         if (buffer_size <= 0) {
416             return -EINVAL;
417         }
418         kfree(dm510_buffer_1->buffer);
419         kfree(dm510_buffer_0->buffer);

```

```

415     dm510_buffer_0->buffer = kmalloc(sizeof(char*) * buffer_size ,
GFP_KERNEL);
416     if (!dm510_buffer_0->buffer){
417         return -ENOMEM;
418     }
419     dm510_buffer_1->buffer = kmalloc(sizeof(char*) * buffer_size ,
GFP_KERNEL);
420     if (!dm510_buffer_1->buffer){
421         return -ENOMEM;
422     }
423
424     dm510_buffer_0->end = dm510_buffer_0->buffer + buffer_size;
425     dm510_buffer_1->end = dm510_buffer_1->buffer + buffer_size;
426
427     dm510_buffer_0->buffer_rp = dm510_buffer_0->buffer;
428     dm510_buffer_0->buffer_wp = dm510_buffer_0->buffer;
429
430     dm510_buffer_1->buffer_rp = dm510_buffer_1->buffer;
431     dm510_buffer_1->buffer_wp = dm510_buffer_1->buffer;
432
433     return buffer_size;
434
435     case IOC.RESETPROC:
436         if (arg <=1) {
437             return -EINVAL;
438         }
439         number_proc = arg;
440         return number_proc;
441
442     default:
443         return -ENOTTY;
444 }
445 return 0;
446 }
447
448 module_init( dm510_init_module );
449 module_exit( dm510_cleanup_module );
450
451 MODULE_AUTHOR( "Michelle Dung Hoang, Lea Fog-Fredsgaard, Danny Rene
Jensen" );
452 MODULE_LICENSE( "GPL" );

```

7.2 iotest.h

```

1  #ifndef IOCTHEADER.h
2  #define IOCTHEADER.h
3
4  #define DM510_IOC_MAGIC 9
5
6  #define IOC.RESETBUFFER _IO(DM510_IOC_MAGIC, 0)
7  #define IOC.RESETPROC _IO(DM510_IOC_MAGIC, 1)
8
9  #define IOC.NUMCASES 2
10
11 #endif

```

7.3 iotest.c


```

1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <fcntl.h>
4 #include <string.h>
5 #include <errno.h>
6 #include <sys/ioctl.h>
7 #include "ioct.h"
8
9
10 int main(int argc, char **argv){
11     char *filename = argv[1];
12     int fd = open(filename, ORDWR);
13     int arg = atoi(argv[3]);
14     int result = 0;
15     char *IO_arg=argv[2];
16     if (!strcmp("buffer size", IO_arg)) {
17         result = ioctl(fd, IOC.RESETBUFFER, arg);
18         printf("changed buffer size to: %d\n", result);
19     } else if (!strcmp("processes", IO_arg)){
20         result = ioctl(fd, IOC.RESETPROC, arg);
21         printf("changed number of max processes to: %d\n", result);
22     } else{
23         printf("input error\n");
24     }
25     close(fd);
26 }

```

7.4 extratests.c

```

1 #include "ioct.h"
2 #include <stdio.h>
3 #include <unistd.h>
4 #include <sys/types.h>
5 #include <stdlib.h>
6 #include <fcntl.h>
7 #include <errno.h>
8 #include <sys/ioctl.h>
9
10
11
12
13
14
15 void writing(){
16     printf("write to buffer:\n");
17     int fd = open("/dev/dm510-0", ORDWR);
18     printf("buffer size changed to %d\n", ioctl(fd, IOC.RESETBUFFER,
19         15));
20
21     //close(fd);
22     //fd = open(filename, ORDWR);
23     char * fillText = "hello";
24     int count = 5, write_result = 0;
25     write_result = write(fd, fillText, count);
26     printf("the result of insert to buffer 0: %d\n", write_result);
27     printf("\n");
28     close(fd);

```

```

28 }
29 }
30
31 void writing1(){
32     printf("write to buffer:\n");
33     int fd1 = open("/dev/dm510-1", ORDWR);
34     //printf("buffer size changed to %d\n", ioctl(fd, IOC.RESETBUFFER,
35         15));
36
37     //close(fd);
38     //fd = open(filename, ORDWR);
39     char * fillText1 = "world";
40     int count = 5, write_result = 0;
41     write_result = write(fd1, fillText1, count);
42     printf("the result of insert to buffer 1: %d\n", write_result);
43
44     printf("\n");
45     close(fd1);
46
47 }
48
49
50
51
52 void reading(){
53     printf("read from buffer0:\n");
54     int fd = open("/dev/dm510-1", ORDWR);
55     int count = 5, read_result = 0;
56     char *text;
57
58     read_result = read(fd, text, count);
59     printf("result from buffer 0 is %d and line: %s\n", read_result,
60         text);
61     printf("\n");
62     close(fd);
63 }
64
65
66 void reading1(){
67     printf("read from buffer1:\n");
68     int fd = open("/dev/dm510-0", ORDWR);
69     int count = 5, read_result = 0;
70     char *text1;
71
72     read_result = read(fd, text1, count);
73     printf("result from buffer 0 is %d and line: %s\n", read_result,
74         text1);
75     printf("\n");
76     close(fd);
77
78 }
79
80
81 void fullBuffers(){

```

```

82 printf("buffer is full:\n");
83 printf("trying to write 15 bytes to buffer twice:\n");
84 int fd = open("/dev/dm510-0", ORDWR);
85 //printf("buffer size changed to %d\n",ioctl(fd, IOC_RESETBUFFER,
    15));
86 //int fd = open(filename, ORDWR);
87 char * fillText = "theBufferIsFull";
88 int count = 15, write_result = 0;
89
90 for (size_t i = 0; i < 2; i++) {
91     write_result = write(fd, fillText, count);
92     if (write_result < 0) {
93         printf("%s \n",strerror(errno));
94     }else{
95         printf("the result of %d insert: %d\n", i, write_result);
96     }
97 }
98 printf("\n");
99 close(fd);
100 }
101
102 void readEmptyBuffers(){
103     printf("read empty buffer:\n");
104     int fd = open("/dev/dm510-0", ORDWR);
105     int count = 15, read_result = 0;
106     char *text;
107     read_result = read(fd, text, count);
108     printf("result is %d\n", read_result);
109     printf("\n");
110     close(fd);
111 }
112
113 void writeNothing(){
114     printf("write nothing:\n");
115     int fd = open("/dev/dm510-0", ORDWR);
116     char * fillText = "";
117     int count = 0, write_result = 0;
118
119     printf("buffer size changed to %d\n",count);
120     write_result = write(fd, fillText, count);
121
122     if (write_result < 0) {
123         printf("%s \n",strerror(errno));
124     }else{
125         printf("result of writing: %d\n", write_result);
126     }
127     printf("\n");
128     close(fd);
129 }
130
131 void readNothing(){
132     printf("read nothing:\n");
133     int fd = open("/dev/dm510-0", ORDWR);
134     int count = 0, read_result = 0;
135     char *text;
136     read_result = read(fd, text, count);
137     if (read_result < 0) {

```

```

138     printf("%s \n",strerror(abs(read_result)));
139 }else{
140     printf("result is %d\n", read_result);
141 }
142 printf("\n");
143 close(fd);
144 }
145
146
147 int main (int argc, char **argv){
148     writing();
149     reading();
150     sleep(1);
151     writing1();
152     reading1();
153     sleep(1);
154     fullBuffers();
155
156     return 0;
157 }

```