

dm510 project 3

Danny Jensen
danje14

Group:
Lea Fog-Fredsgaard
Michelle Dung Hoang

April 2018

Contents

1	Introduction	1
2	Design	1
3	Implementation	2
4	Test	5
5	Conclusion	6
6	Appendix	7
6.1	lfs.c	7
6.2	makefile	15

1 Introduction

In this project a file system must be implemented. The user of the file system must be able to create and navigate through the newly created folders. The user should also be able to create files that the user can then write and read from using *echo* and *cat*. All these folders and file should be stored in an inode.

2 Design

The first design choices was how the inode structure should look. Since it was a requirement that an access time and a modified time was present this was the first thing to look at. The the choice was made that the inode also had to have its path as its name so that it would be easier to search for a given inode. every inode whould then also have a "unique" number to help locate the place where it would be saved onto the disk.

Sadly the indirect pointers was not implemented because of errors. But the inode holds pointers to all its children which makes it easy to navigate through every layer.

The inode of-cause also need some place to hold data, this was first tried with and string array but the idea was dropped because it needed a fixed length and that also meant that the inode would use more space than necessary. Instead a void pointer is used that points to the place in memory the data is put and an integer for the length is added.

Since the structure must be inside blocks that must be inside segments. A block structure is needed, this block should hold the inode and some integer so as to identify if the inode is usable or not.

Since it was impossible to get the segment part working it was dropped and the blocks are simply saved to the disk.

To retrieve the inode, the first plan was to let a function run through the linked list structure but that was dismissed as it would be much harder than simply have an array with every inode in that is in the memory.

To save everything in a file, the first idea was to save every change to a file as soon as it was made, but this would make the process run slower because of saves and retrieves all the time, that is why a simple void variable is chosen. This variable is updated with all new info and should be retrieved when the process starts and saved to a file when the process is shut down, sadly there was not enough time to implement this part.

Another design decision was that it would not be possible to delete a folder with contents in it, this design decision was taken because the *unlink* function should not be implemented and that meant that there was no way of deleting a file.

3 Implementation

As mentioned in the Design chapter, the inode has to have ways to store the time of last access and modified, these variables are saved as *time_t*. The inode also needs a type to tell if it is a folder or a file, for this a simple int variable is used, if the variable is 0 it is a folder and if it is 1 the inode is a file. The *inode_direct* variable is a set of pointers to all the direct children of the inode, and the *inode_path* is the inode's name and its path as a string.

The *current_inode_amount* variable is the one that holds the number of variables. No two inodes should have the same number. This variable is incremented each time a new inode is created.

The *lfs_init* function simply initializes all the variables such as *disk* and *array_inode* and allocate space for it in memory. *lfs_init* also creates the root inode that is the root folder, and puts it into *array_inode* at place 0.

There are three different functions to find inodes, *find_inode*, *find_inode_dir* and *find_inode_file* they pretty much do the same. They look through *array_inode* to find the inode with the given path and returns it. The difference is that *find_inode* does not care what type of inode it is, either folder or file, where *find_inode_dir* finds only a folder with that path and *find_inode_file* only finds the file with that path.

```
1 if ((strcmp(array_inode[i] -> node -> inode_path, path) == 0) && (
    array_inode[i] -> node -> inode_type == 0)){
```

These three functions could have been easily put into one function but the time did not allow for optimization of the code.

To find the parent of a file or folder a simple function named **find_inode_parent* was created. It copies the path to another string variable, this is done because *path* is a *const char* type and can not be changed. On this new copy the *dirname* function is used, it is a smart function that removes everything after the last / in a string. This new string is then the parent path and can simply be put into a *find_inode* function to return the parent of a path.

```
1 inode *find_inode_parent(const char *path){
2     char *path_copy;
3     path_copy = strdup(path);
4     return find_inode(dirname(path_copy));
5 }
```

lfs_getattr simply copies the data from the inode, found by the path, over into the *stbuf* with the data that was given. Just to ensure that nothing is wrong *res* is set to *ENOENT* if there is no inode with that path or that something went wrong and the type can't be right.

lfs_readdir first creates an inode to put the hopefully found inode into, with the given path. To find this inode *find_inode_dir* is used because as the name states it needs to return a folder. When the right inode is found, every path

of every child is then given to the *filler* function to return to the user. This is easily done because the inode has pointers to all its direct children, and an int variable that holds the number of children. Just like *lfs_getattr*, *lfs_readdir* will also return *ENOENT* if there is no inode with that path.

```
1 for(int i = 0; i < inod -> inode_sub_num; i++){
2     filler(buf, basename(((inode*)inod -> inode_direct[i]) ->
3         inode_path), NULL, 0);
}
```

To simplify *lfs_mknod* and *lfs_mkdir* they were simply put into one function named *lfs_mk* that takes the path for the new inode and the type in form of an int

```
1 int lfs_mkdir(const char *path, mode_t mode) {
2     int ret = lfs_mk(path, 0);
3     return ret;
4 }
```

The *lfs_mknod* and *lfs_mkdir* could easily have been one liners there is no real meaning in creating an int variable to return, the return value from *lfs_mk* could just have been returned, but at the time this gave sense and as stated before sadly there was not enough time to optimize the code.

lfs_mk first finds the parent of the inode that is to be created, this is done by using *find_inode_parent*. then *lfs_mk* starts to allocate space for the new inode. There is also allocated space for *inode_data* which shall hold the data if it is a file, this should actually have been inside an if statement so as to not use unnecessary ram for a folder inode which has no use for it. The size is set to 20 times the size of a void pointer, this was done to test the ability to save strings inside files, and should have been changed, but it was down prioritized over getting other things working.

Now that the new inode with the given path is created, the parent inode needs to get a pointer in its *inode_direct* and of course its *inode_sub_num* needs to be incremented, here a wierd thing happened if a simple

```
1 parent -> inode_sub_num++;
```

was used first but then the the process froze, so

```
1 parent -> inode_sub_num = parent -> inode_sub_num+1;
```

was used and it fixed the error.

Now *lfs_rmdir* was a bit tricky. First the parent is found for later use. Then the actual inode that is to be deleted is found, and since it needs to be a folder is *find_inode_dir* used.

If the folder has no children then it can be deleted, as stated in the design chapter this is so because if there was a file inside the folder to be deleted it should also be deleted but there is no way of deleting a file in the code.

If the folder can be deleted, the pointer in the parent is first deleted. This is done by setting the pointer to *NULL* and then simply run through every other pointer and placing it in the previous place in the array.

```

1 for(int j = i; j < parent -> inode_sub_num-1; j++){
2     parent -> inode_direct[j] = parent -> inode_direct[j+1];
3 }

```

After this the last pointer is set to *NULL* and *inode_sub_num* of the parent is reduced by one.

The next part has some errors in it.

```

1 for(int i = 0; i < current_inode_amount; i++){
2     if(strcmp(array_inode[i] -> node -> inode_path, inod ->
3         inode_path) == 0){
4         array_inode[i] = NULL;
5         for(int j = i; j < current_inode_amount; j++){
6             array_inode[j] = array_inode[j+1];
7         }
8         array_inode[current_inode_amount] = NULL;
9         current_inode_amount--;
10        break;
11    }
12 }

```

In an effort to try and clean the data a *cleanup* function was meant to be added and this is the remnants of this function that was forgotten in the code. If *current_inode_amount* is reduced and then a new node is added the data could be stored over an inode that is still in use inside the disk. This will lead to nodes being lost and that is not a good thing. The reduction on the *array_inode* is okay because the "unique" number of the inodes are stored in *inode_number* of each inode.

This is a bit confusing so lets take a small example, every node is a folder:

3 nodes are created with *inode_number*, 1,2 and 3. Now *current_inode_amount* is 3. node 2 is then deleted and *current_inode_amount* is reduced. Then a new node is created this means that it get 3 in its *inode_number*, this means we have 1, 3 and 3. And since the *inode_number* is used to get the pointer to where the inode is stored this means that both inodes are stored the same place which is not possible. This means that the first inode with *inode_number* that is 3 is overwritten.

lfs_read is pretty simple, it just finds the inode with the help of *find_inode_file* and copies everything inside the *inode_data* to the buffer and returns the number of bytes stored in the *inode_size* of the inode.

lfs_write finds the inode using *find_inode_file* copies the string over into *inode_data*. It then finds the length of the string and puts it into the *inode_size* of the inode.

```

1 memcpy(inod -> inode_data, buf, strlen(buf));
2 inod -> inode_size = strlen(buf);

```

When is done with putting the string into the inode, the inode is then put into the disk with the new information. The inode is then put into a block and *array_inode* is updated. This part gave some problems because it did not make so much sense why *array_inode* should be updated, it only has pointers

to the inodes inside it. But an error was returned every time the inode was read.

lfs_utime simply retrieves the inode and updates the *inode_access* and *inode_modified* time with the time given from the *utb*.

4 Test

The test was video captured so it will be refereed to with time stamps here instead of showing pictures.

First the file system is mounted, this happens at 00.05. And then go into the newly mounted file system.

An *ls* command is run to show that the folder is empty at 00.19. And right after that a folder is created named *dir1*

At 00.29 another *ls* is run to show that the *mkdir* and the *ls* functions actually work, and as shown a directory was created with the name *dir1*.

Next a file is created with the name *file1* using the *touch* command, this happens at 00.34

Just to show that *file1* was created a *ls* command was run again. And as it shows *file1* was created.

at 00.46 the *cd* command was checked by going into the directory by the name *dir1*. And as the path shows, *dir1* was successfully entered. Just to show that this directory is also empty a *ls* command is run in this directory too.

Now to navigate back to the parent folder of *dir1* the *cd ..* is used, this happens at 0.52. And this is also successful as shown by the path to the right.

A new *ls* command is used to show that the folder and file both still are present in the root directory, this happens at 0.65.

A new directory is made using *mkdir* by the name *dir2* and then *cd* is used to navigate into this folder, at 1.00. Then at 1.17 another file is created using the *touch* command, this new file is named *file2*

at 1.44 *file2* is written to using the *echo* command, the string that was written was "*hello from file2*".

Now to see if the string added to *file2* really was added and saved, a *cat* command is used on *file2* at 1.50. And yes the string was added and saved as it is printed out to the terminal.

Now go back to the root and try to delete folders. To delete a folder the *rmdir* command is used, and the target directory to remove is *dir1* this happens at 2.06. a *ls* command is run to show that *dir1* was removed from the root directory.

dir1 was empty when removed so what happens if the directory that is to be deleted not empty is. This test is performed at 2.15. *dir2* is not empty as it holds *file2* but the *rmdir* command is still performed on *dir2*. As shown in the output the folder *dir2* is not empty and *rmdir* failed to remove it.

Now to test the *ls* function a bit *ls* is performed with *dir2* as a path while standing in the root directory, this happens at 2.33. As shown *dir2* holds *file2*,

so this also works like planned.

Next to test if *cat* can do the same. At 2.43 *cat* is given the path *dir2/file2* while also still in the root directory. "*hello from file2*" was printed out to the terminal and this was also a huge success.

5 Conclusion

The file system is able to create and it is possible to navigate through the directories. It is also possible to create, write to and read from a file.

Sadly it is not possible to shut the file system down and start it back up again without loosing all the data, but this should not be that hard to do since everything is in a *disk* variable.

The segmentation part is not completed either but if the user does not use too long strings this is no problem. And the code could use some cleanup from other things that was created in an attempt to fix other errors.

But all in all the code runs and it is possible to use it as a rudimentary file system.

6 Appendix

6.1 lfs.c

```
1 #include <fuse.h>
2 #include <errno.h>
3 #include <string.h>
4 #include <stdio.h>
5 #include <stdlib.h>
6 #include <time.h>
7 #include <fcntl.h>
8 #include <libgen.h>
9 #include <sys/types.h>
10 #include <unistd.h>
11
12 #define SEGMENTS 4      /* segments containing the blocks, and
13                          array_inode */
14 #define BLOCK 4         /* Blocks containing the data */
15 #define BLOCK_SIZE 256
16 #define SEGMENT_SIZE 1024
17
18 typedef struct inode{
19     ino_t inode_number;      /* id of the node */
20     char *inode_path;        /* path to where the file or directory
21                              is placed on the computer*/
22     int inode_type;          /* 0 directory, 1 file */
23     size_t inode_sub_num;    /* size of sub inodes in current inode
24                              */
25     time_t inode_access;     /* time of last access */
26     time_t inode_modified;   /* time of last modification */
27     void **inode_direct;     /* pointer to all sub files or
28                              directories */
29     //int **inode_indirect;   /* pointer to pointer to datablocks */
30     size_t inode_size;       /* size of data in files */
31     void *inode_data;        /* data */
32 } inode;
33
34 typedef struct block {
35     struct inode *node;
36     int in_use;
37 } block;
38
39 inode *find_inode(const char* path);
40 inode *find_inode_dir(const char* path);
41 inode *find_inode_file(const char* path);
42 inode *find_inode_parent(const char* path);
43 int lfs_getattr(const char *, struct stat *);
44 int lfs_readdir(const char *, void *, fuse_fill_dir_t , off_t ,
45                struct fuse_file_info *);
46 int lfs_mk(const char *path, int type);
47 int lfs_mknod(const char* path, mode_t mode, dev_t rdev);
48 int lfs_mkdir(const char *path, mode_t mode);
49 int lfs_rmdir(const char *path);
```

```

49 int lfs_open(const char *, struct fuse_file_info *);
50 int lfs_read(const char *, char *, size_t, off_t, struct
    fuse_file_info *);
51 int lfs_release(const char *path, struct fuse_file_info *fi);
52 int lfs_write(const char *path, const char *buf, size_t size, off_t
    offset, struct fuse_file_info *fi );
53 int lfs_utime (const char *path, struct utimbuf *utb);
54
55
56
57
58 block **array_inode;          /* pointer to the inodes */
59 int current_inode_amount = 0; /* current amount of inodes,
    increased when a new node is made */
60 void *disk;                  /* disk containing the segments
    */
61 int segment;
62 //int number_of_inodes = 0;
63
64 int lfs_init(void){
65     printf("——init——\n");
66     inode *inod;
67     inod = malloc(sizeof(inode)); /* allocate space for the inode
    */
68     if(inod == NULL){
69         return -ENOMEM;
70     }
71     inod -> inode_path = calloc(200, sizeof(char)); /* allocate space
    for the inode path */
72     if(inod -> inode_path == NULL){
73         return -ENOMEM;
74     }
75     char *root_path = "/";
76     /* initializing the rootnode */
77     inod -> inode_number = 0; /* Root node */
78     strcpy(inod -> inode_path, root_path); /* initialize the root
    path to "/" */
79     inod -> inode_type = 0;
80     inod -> inode_sub_num = 0;
81     inod -> inode_size = 0;
82     inod -> inode_access = time(NULL);
83     inod -> inode_modified = time(NULL);
84     inod -> inode_direct = malloc(15*sizeof(inode));
85     //printf("hello1\n");
86     array_inode = malloc(1024 * sizeof(inode));
87     if(array_inode == NULL){
88         free(inod);
89         return -ENOMEM;
90     }
91
92     disk = malloc(SEGMENTS * SEGMENT_SIZE);
93     if(disk == NULL){
94         free(inod);
95         free(array_inode);
96         return -ENOMEM;
97     }
98     /*

```

```

99     segment = malloc(BLOCK * BLOCK_SIZE);
100     if(segment == NULL){
101         free(inod);
102         free(array_inode);
103         free(disk);
104         return -ENOMEM;
105     }
106 */
107 //printf("hello2\n");
108 block *bl;
109 bl = malloc(sizeof(block));
110
111 if(bl == NULL){
112     free(inod);
113     free(array_inode);
114     free(disk);
115     //free(segment);
116     return -ENOMEM;
117 }
118
119 bl -> node = inod;
120 bl -> in_use = 1;
121
122 array_inode[0] = bl; /* places the rootnode in the 0'th place */
123 current_inode_amount++;
124 memcpy((disk + (current_inode_amount-1 * BLOCK_SIZE)), inod,
        sizeof(inode)); /* copies from inod to the next available space
        on the disc */
125
126 //printf("GOODBYE from init\n");
127 return 0;
128 }
129
130 /* find inodes */
131 inode *find_inode(const char* path){
132     printf("——find_inode——\n");
133     for(int i = 0 ; i < current_inode_amount ; i++) {
134         printf("        find_inode: %i out of %i iterations and path: %s\n", i, current_inode_amount, path);
135
136         if (strcmp(array_inode[i] -> node -> inode_path, path) == 0){
137             /* if it finds the path */
138             printf("        the path of the inode is found: %s\n", array_inode[i] -> node -> inode_path);
139             inode *inod;
140             //inod = malloc(sizeof(inode));
141             inod = array_inode[i] -> node;
142             return inod; /* returns the inode*/
143         }
144     }
145     printf("no inode found \n");
146     return NULL;
147 }
148
149 /* find directory inodes */
150 inode *find_inode_dir(const char* path){

```

```

150 printf("——find_inode_dir——\n");
151 for(int i = 0 ; i < current_inode_amount ; i++) {
152     printf("        find_inode_dir: %i out of %i iterations and path
153         %s\n", i, current_inode_amount, path);
154
155     /* if it finds the path and it is a directory */
156     if ((strcmp(array_inode[i] -> node -> inode_path, path) == 0)
157         && (array_inode[i] -> node -> inode_type == 0)){
158         printf("        the path of the directory is found: %s\n",
159             array_inode[i] -> node -> inode_path);
160         inode *inod;
161         inod = malloc(sizeof(inode));
162         inod = array_inode[i] -> node;
163         return inod;                                /* returns the
164         inode*/
165     }
166 }
167 printf("no find_inode_dir found \n");
168 return NULL;
169 }
170
171 /* find file inodes */
172 inode *find_inode_file(const char *path){
173     //inode *inod;
174     printf("——find_inode_file——\n");
175     for(int i = 0 ; i < current_inode_amount ; i++) {
176         printf("        find_inode_file: %i out of %i iterations and path
177         %s\n", i, current_inode_amount, path);
178
179         /* if it finds the path and it is a file */
180         if ((strcmp(array_inode[i] -> node -> inode_path, path) == 0)
181             && (array_inode[i] -> node -> inode_type == 1)){
182             printf("        the path of the file is found: %s\n",
183                 array_inode[i] -> node -> inode_path);
184             return (array_inode[i] -> node);
185             /* returns the inode*/
186         }
187     }
188     printf("no find_inode_file found \n");
189     return NULL;
190 }
191
192 /* find parent inode */
193 inode *find_inode_parent(const char *path){
194     printf("——find_parent——\n");
195     char *path_copy;
196     path_copy = strdup(path);
197     return find_inode(dirname(path_copy));    /* copies the path up
198     to, but not including, the final '/' */
199 }
200
201 /* */
202 int lfs_getattr( const char *path, struct stat *stbuf ) {
203     printf("——getattr——(path: %s)\n", path);
204     int res = 0;

```

```

198 //printf("getattr: (path=%s)\n", path);
199
200 memset(stbuf, 0, sizeof(struct stat)); /* initialize stbuf with 0
    */
201 inode *inod;
202 //if(inod == NULL){
203 //    return -ENOENT; /* no such file og
        directory */
204 //}
205 //printf("getattr3\n");
206 inod = find_inode(path); /* finds the inode */
207 if(inod == NULL){
208     return -ENOENT;
209 }
210 stbuf->st_ino = current_inode_amount;
211 stbuf->st_size = inod->inode_size;
212 stbuf->st_atime = time(NULL);
213 stbuf->st_mtime = inod->inode_modified;
214 stbuf->st_ctime = inod->inode_access;
215 if( inod->inode_type == 0){ /* directory */
216     stbuf->st_mode = S_IFDIR | 0755;
217     stbuf->st_nlink = 2;
218 }else if(inod->inode_type == 1){ /* file */
219     stbuf->st_mode = S_IFREG | 0777;
220     stbuf->st_nlink = 1;
221 }else /* file */
222     res = -ENOENT; /* no such file og directory
        */
223 //printf("getattr end\n");
224 printf("        last access time: %li \n", inod->inode_access);
225 printf("        last modified time: %li\n\n", inod->
        inode_modified);
226 return res;
227 }
228
229
230 /* display all content of a directory */
231 int lfs_readdir(const char *path, void *buf, fuse_fill_dir_t filler
        , off_t offset, struct fuse_file_info *fi ) {
232     printf("——readdir——\n");
233     (void) offset;
234     (void) fi;
235
236     inode *inod = malloc(sizeof(inode));
237     if(inod == NULL){
238         return -ENOENT; /* no such file og directory */
239     }
240
241     filler(buf, ".", NULL, 0);
242     filler(buf, "..", NULL, 0);
243     inod = find_inode_dir(path); /* Finds the path given to
        lfs_readdir, in the current directory */
244     /*
245     if(strcmp(path, "/") != 0){
246         return -ENOENT;
247     }*/
248

```

```

249     for(int i = 0; i < inod -> inode_sub_num; i++){
250         filler(buf, basename(((inode*)inod -> inode_direct[i]) ->
            inode_path), NULL, 0); /* Filler displays the content of the
            buffer with all inodes in the current directory */
251     }
252     printf("——readdir end——\n");
253     //free(inod);
254     return 0;
255 }
256
257
258 int lfs_mk(const char *path, int type) {
259     printf("——mk——\n");
260     inode *parent;
261     parent = malloc(sizeof(inode));
262     if(parent == NULL){
263         return -ENOMEM;
264     }
265     printf("mk2\n");
266     parent = find_inode_parent(path);
267
268     inode *inod; /* file or directory */
269     inod = malloc(sizeof(inode)); /* allocate space for the inode
        */
270     if(inod == NULL){
271         return -ENOMEM;
272     }
273     printf("mk3\n");
274     inod -> inode_path = malloc(sizeof(char)); /* allocate space for
        the inode path */
275     if(inod -> inode_path == NULL){
276         return -ENOMEM;
277     }
278     printf("mk4\n");
279     inod -> inode_number = current_inode_amount; /* current
        node */
280     strcpy(inod -> inode_path, path); /* initialize
        the path */
281     inod -> inode_type = type;
282     inod -> inode_sub_num = 0;
283     inod -> inode_size = 0;
284     inod -> inode_access = time(NULL);
285     inod -> inode_modified = time(NULL);
286     inod -> inode_direct = malloc(15*sizeof(inode));
287     inod -> inode_data = malloc(20*sizeof(void*));
288     block *bl;
289     bl = malloc(sizeof(block));
290     if(bl == NULL){
291         return -ENOMEM;
292     }
293     bl -> node = inod;
294     bl -> in_use = 1;
295     array_inode[current_inode_amount] = bl; /* places the current
        block with the inode in the array */
296     current_inode_amount++;
297     memcpy((disk + (current_inode_amount * BLOCK_SIZE)), inod, sizeof
        (inode)); /* copies from inode to the next available space on

```

```

    the disc */
298 parent -> inode_direct[parent -> inode_sub_num] = inod; /* points
    to the new inode */
299 parent -> inode_sub_num = parent -> inode_sub_num+1; /*
    increment size of the inode */
300 return 0;
301 }
302 /* uses lfs_mk to create a file */
303 int lfs_mknod(const char* path, mode_t mode, dev_t rdev) {
304     printf("——mknod——\n");
305     int ret = lfs_mk(path, 1);
306     return ret;
307 }
308
309 /* uses lfs_mk to create a directory */
310 int lfs_mkdir(const char *path, mode_t mode) {
311     printf("——mkdir——\n");
312     int ret = lfs_mk(path, 0);
313     return ret;
314 }
315
316 int lfs_rmdir(const char *path) {
317     printf("——rmdir——\n");
318     inode *parent;
319     parent = malloc(sizeof(inode));
320     if(parent == NULL){
321         return -ENOMEM;
322     }
323     parent = find_inode_parent(path);
324
325     inode *inod;
326     inod = malloc(sizeof(inode)); /* directory to be deleted */
327     /* allocate space for the inode */
328     if(inod == NULL){
329         return -ENOMEM;
330     }
331     inod = find_inode_dir(path); /* the path to the directory to
    be deleted */
332
333     if(inod -> inode_sub_num == 0){ /* is there sub
    directories */
334         for(int i = 0; i < parent -> inode_sub_num; i++){
335             //inod = parent -> inode_direct[i]; /* finds the pointer to
    the parent directory */
336             if(strcmp(((inode*)parent -> inode_direct[i]) -> inode_path,
    inod -> inode_path) == 0){ /* finds the right path in the
    parent */
337                 parent -> inode_direct[i] = NULL;
338                 for(int j = i; j < parent -> inode_sub_num-1; j++){
339                     parent -> inode_direct[j] = parent -> inode_direct[j+1];
340                     /* when found, shift everything one place to the left */
341                 }
342                 parent -> inode_direct[parent -> inode_sub_num-1] = NULL;
343                 parent -> inode_sub_num--;
344                 break;
345             }
346         }
347     }
348 }

```

```

345     for(int i = 0; i < current_inode_amount; i++){
346         if(strcmp(array_inode[i] -> node -> inode_path, inode ->
347             inode_path) == 0){ /* finds the right path in the in the inode
array */
348             array_inode[i] = NULL;
349             for(int j = i; j < current_inode_amount; j++){
350                 array_inode[j] = array_inode[j+1]; /* when found, shift
everything one place to the left */
351             }
352             array_inode[current_inode_amount] = NULL;
353             current_inode_amount--;
354             break;
355         }
356     }
357     return 0;
358 }
359
360 return -ENOTEMPTY;
361 }
362
363 //Permission
364 int lfs_open(const char *path, struct fuse_file_info *fi) {
365     printf("——Open—— path: %s\n", path);
366     return 0;
367 }
368
369 int lfs_read(const char *path, char *buf, size_t size, off_t offset
, struct fuse_file_info *fi) {
370     printf("——Read—— path: %s\n", path);
371     inode *inode;
372     inode = find_inode_file(path);
373     memcpy(buf, inode -> inode_data, inode -> inode_size);
374     printf("——read end——\n");
375     return inode -> inode_size;
376 }
377
378 int lfs_release(const char *path, struct fuse_file_info *fi) {
379     printf("——Release—— path: %s\n", path);
380     return 0;
381 }
382
383 int lfs_write(const char *path, const char *buf, size_t size, off_t
offset, struct fuse_file_info *fi) {
384     printf("——Write—— path: %s\n", path);
385     inode *inode;
386     inode = find_inode_file(path);
387     memcpy(inode -> inode_data, buf, strlen(buf));
388     inode -> inode_size = strlen(buf);
389     memcpy((disk + ((inode -> inode_number) * BLOCK_SIZE)), inode,
sizeof(inode));
390     array_inode[inode -> inode_number] = NULL;
391     block *bl = malloc(sizeof(block));
392     bl -> node = inode;
393     bl -> in_use = 1;
394     array_inode[inode -> inode_number] = bl;
395

```



```

396     printf("——write end——\n" );
397     return size;
398 }
399
400 int lfs_utime (const char *path, struct utimbuf *utb) {
401     printf("——Utime——, path: %s\n", path);
402
403     inode *inode;
404     inode = find_inode(path);
405     inode -> inode_access = utb -> actime;
406     inode -> inode_modified = utb -> modtime;
407
408     return 0;
409 }
410
411 static struct fuse_operations lfs_oper = {
412     .getattr  = lfs_getattr ,
413     .readdir  = lfs_readdir ,
414     .mknod   = lfs_mknod ,
415     .mkdir   = lfs_mkdir ,
416     .unlink  = NULL, //not needed
417     .rmdir   = lfs_rmdir ,
418     .truncate = NULL,
419     .open    = lfs_open ,
420     .read    = lfs_read ,
421     .release  = lfs_release ,
422     .write   = lfs_write ,
423     .rename  = NULL, //not needed
424     .utime   = lfs_utime
425 };
426
427 int main( int argc , char *argv[] ) {
428     lfs_init();
429     fuse_main( argc , argv , &lfs_oper );
430     return 0;
431 }

```

6.2 makefile

```

1 GCC = gcc
2 SOURCES = lfs.c
3 OBJS := $(patsubst %.c,%.o,$(SOURCES))
4 CFLAGS = -O2 -Wall -D_FILE_OFFSET_BITS=64 -DFUSE_USE_VERSION=25
5
6 .PHONY: lfs
7
8 ##
9 # Libs
10 ##
11 LIBS := fuse
12 LIBS := $(addprefix -l,$(LIBS))
13
14 all: lfs
15
16 %.o: %.c
17     $(GCC) $(CFLAGS) -c -o $@ $<
18

```

```
19 lfs: $(OBS)
20   $(GCC) $(OBS) $(LIBS) $(CFLAGS) -o lfs
21
22 clean:
23   rm -f $(OBS) lfs
24
25 mount:
26   mkdir lfs-mountpoint1
27   chmod 755 lfs-mountpoint1/
28   ./lfs -f lfs-mountpoint1/
29
30 unmount:
31   fusermount -u lfs-mountpoint1/
32   rm -rf lfs-mountpoint1
```